


PaSe: An Extensible and Inspectable DSL for Micro-Animations

Ruben P. Pieters ^[0000-0003-0537-9403] and Tom
Schrijvers^[0000-0001-8771-5559]

KU Leuven, 3001 Leuven, Belgium
{ruben.pieters, tom.schrijvers}@cs.kuleuven.be

First author is a *research student*. Article category: *Project Article*.

Abstract. This paper presents PaSe, an extensible and inspectable DSL embedded in Haskell for expressing micro-animations. PaSe builds animations in compositional fashion, using parallel and sequential animations as basic building blocks. This differs from typical animation libraries which mostly focus on sequential composition and utilize callbacks and implicit effects for their expressivity. To provide similar flexibility to other animation libraries, PaSe features extensibility of operations and inspectability of animations. We present the features of PaSe with a to-do list application, discuss the PaSe implementation, and argue that the callback style of extensibility is detrimental for correctly integrating inspectability. To illustrate this, we contrast with the GreenSock Animation Platform, a professional-grade and widely used JavaScript animation library.

1 Introduction

Because of their ability to structure effectful code in a pure functional codebase, monads quickly became ubiquitous in functional programming [20]. They have since seen wide use in Haskell Domain Specific Languages (DSLs). However, the choice for a monadic DSL implies certain trade-offs. The obvious advantage of monadic DSLs is their expressivity, but there are also drawbacks. The main loss is that of *inspectability*, as monadic computations can only be inspected up to the next action. Techniques such as applicative functors [16], arrows [9], or selective applicative functors [18] choose the other side of the trade-off: they increase the inspection capabilities by reducing the expressivity compared to monads.

This paper develops a DSL embedded in Haskell for defining micro-animations, called PaSe¹. PaSe employs a technique which alleviates some aspects of the trade-off between expressivity and inspectability. The expressivity of control flow is restricted by means of type classes, inspired by the MTL style originally introduced by Liang *et al.* [14]. The MTL style is an open encoding which allows extensions to the syntax of the DSL. Instantiating the abstract animation definitions with, for example, the **Const** functor provides inspectability. Expressivity

¹ Pronounced *pace* (peis), the name is derived from *Parallel* and *Sequential*.

can be increased, while preserving inspectability, by adding new control flow constructs to the DSL and providing a corresponding instance for inspection.

Micro-animations are short animations displayed when users interact with an application, for example an animated transition between two screens. When used appropriately, they aid the user in understanding evolving states of the application [1,7,8]. Examples can be found in almost every software application: window managers animate window minimization, menus in mobile applications pop in gradually, browsers highlight newly selected tabs with an animation, etc.

PaSe provides the features expected of animation libraries by building them with recent ideas from functional programming. Our contributions are as follows:

- We develop PaSe, which supports arbitrary composition of animations and inspectability. Animation libraries, such as the GreenSock Animation Platform (GSAP)², typically use callbacks as a means of extensibility/expressivity; this is detrimental to inspectability. We show an example resulting in unexpected behaviour and how PaSe correctly handles it.
- PaSe is an *extensible* DSL: the syntax can be extended with new operations. The animations use case is novel for approaches to extensibility.
- PaSe supports *inspectability*: extracting information from computations before running it. Inspectability is present in specific computation classes, such as free applicatives [2]. But, it is novel to combine it with extensibility.
- PaSe supports arbitrary nesting of parallel and sequential animations which correctly interacts with inspectability. Such parallel components exist already, see for example Ren'Py³, React Native Animations⁴ or Qt Animations⁵. Yet, general-purpose animation libraries lack them. Also, we correctly support the interaction with inspectability.
- We implemented various examples⁶: a to-do list application, a communication story example, a game-like demo application and a Pac-Man game. We combined PaSe with both `gloss`⁷ and the Haskell SDL bindings⁸ as graphics backend. This paper uses the to-do list as motivating application and compares the development of the Pac-Man application, developed in both Haskell with PaSe and in TypeScript with GSAP.

2 Motivation

We present a to-do list application to showcase the functionality of PaSe.

² <https://greensock.com>

³ <https://www.renpy.org/doc/html/atl.html#parallel-statement>

⁴ <https://facebook.github.io/react-native/docs/animated#parallel>

⁵ <https://doc.qt.io/qt-5/animation.html>

⁶ <https://github.com/rubenpieters/PaSe-hs/tree/master/PaSe-examples>

⁷ <https://hackage.haskell.org/package/gloss>

⁸ <https://hackage.haskell.org/package/sdl2>

71 2.1 Running Example

72 Our application has two screens: a main screen and a menu screen. The main
 73 screen contains a navigation bar and three items. An overview of the application
 74 is given in Figure 1. These screenshots are captured from the application built
 75 by combining PaSe with `gloss` as graphics backend.




Fig. 1: Overview of the to-do list application.

76 In this application, various user actions are accompanied with an animation.
 77 We list these actions below. Some animations are shown in Figure 2.

- 78 – The user marks items as *(not)* done by clicking them. The checkmark icon
 79 changes shape and color to display its status change.
- 80 – The user filters items by their status with the navigation bar buttons. The
 81 leftmost shows all items, the middle shows all completed items, and the
 82 rightmost shows all unfinished items. The navigation bar underline and to-
 83 do items itself change shape to indicate the new selection.
- 84 – The menu screen shows/hides itself after clicking the menu icon (`≡`). The
 85 menu expands inwards from the left, to indicate the application state changes.

86 2.2 Composing Animations

87 Animations are built in a compositional fashion. When creating an animation, we
 88 decompose it into smaller elements. For example, the `menuIntro` animation both
 89 introduces the menu screen and fades out the background. Thus, it is composed
 90 of two basic animations `menuSlideIn` and `appFadeOut` in parallel. The next
 91 sections explain how to construct such basic and composed animations.

92 **Basic Animations** Basic animations change the property of an element over
 93 a period of time. The `linearTo` function has three inputs: a lens targeting the
 94 property, the duration, and the target value for this property. This results in a
 95 linear change from the current value to its target, hence the name. The duration
 96 is specified with `For` while the target value is specified with `To`.

97 To animate the navigation bar underline, we reduce the width of the leftmost
 98 underline for 0.25 seconds and increase the width of the middle underline for 0.25
 99 seconds. These animations are expressed in respectively `line1Out` and `line2In`
 100 below, and visualized in Figure 3.

```
line1Out = linearTo (navbar . underline1 . width) (For 0.25) (To 0)
line2In = linearTo (navbar . underline2 . width) (For 0.25) (To 28)
```




Fig. 2: Micro-Animations in the to-do list application.

101 *Note on Lenses* We use lens notation `x . y . z` to target `z` inside a nested
 102 structure `{ x: { y: { z: T } } }`. This type of lenses was conceived by van
 103 Laarhoven [13], and later packaged into various Haskell libraries, such as `lens`⁹.

104 The `menuSlideIn` and `appFadeOut` animations are other examples. For the
 105 former, we increase the width of the menu over a duration of 0.5 seconds, and
 106 for the latter we increase the opacity of the obscuring box, determined by `alpha`,
 107 over a duration of 0.5 seconds. These animations are visualized in Figure 3.





```
menuSlideIn = linearTo (menu . width) (For 0.5) (To 75)
appFadeOut  = linearTo (obscuringBox . alpha) (For 0.5) (To 0.65)
```

108 **Composed Animations** A composed animation combines several other ani-
 109 mations into a new one. We can do this either in *sequence* or in *parallel*.

110 We create `selectBtn2` by combining `line1Out` and `line2In` with `sequential`.
 111 This constructs a new animation which first plays `line1Out`, and once it is fin-
 112 ished plays `line2In`.

```
selectBtn2Anim = line1Out 'sequential' line2In
```

⁹ <https://hackage.haskell.org/package/lens>

(a) The `line1Out` animation.(b) The `line2In` animation.(c) The `menuSlideIn` animation.(d) The `appFadeOut` animation.Fig. 3: Basic `linearTo` animations.

113 To obtain `menuIntro`, we combine both `menuSlideIn` and `appFadeOut` with
 114 `parallel`. This constructs a new animation which plays both `menuSlideIn` and
 115 `appFadeOut` at the same time.

```
menuIntro = menuSlideIn 'parallel' appFadeOut
```

116 Both of these animations are visualized in Figure 4.

117 3 Extensibility, Inspectability and Expressiveness

118 The features in Section 2 form the basis of PaSe. To provide support for addi-
 119 tional features present in other animation libraries, we design PaSe to be extensi-
 120 ble and inspectable. This means that PaSe can be extended with new operations
 121 and information can be derived from inspecting specified animations. To support
 122 arbitrary expressiveness in combination with those features, we also emphasize
 123 the possibility to extend PaSe with new combinators.



(a) The `selectBtn2` animation.(b) The `menuIntro` animation.

Fig. 4: All of the defined composed animations.

124 3.1 Extensibility

125 The `linearTo` operation and the `sequential` and `parallel` combinators form
 126 the basis for expressing a variety of animations. However, there are situations
 127 which require other primitives to express desired animations. For example, GSAP
 128 provides a primitive to morph one shape into another.

129 An example in the to-do list app is `checkIcon`, part of `markAsDone`, where
 130 we want to set the color of the checkmark to a new value. We define a custom
 131 `set` operation and embed it inside a PaSe animation. In this animation we use
 132 Haskell's `do`-notation to specify sequential animations.

```
checkIcon = do ...; set (checkmark . color) green; ...
```

133 3.2 Inspectability

134 PaSe is inspectable, meaning that we can derive properties of expressed compu-
 135 tations by *inspecting them rather than running them*. For example, we want to
 136 know the duration of `menuIntro` without actually running it and keeping track
 137 of the time. The `duration` function calculates the duration by inspecting the
 138 animation. Passing it `menuIntro` gives a duration of `0.5` seconds, which is indeed
 139 the duration of two `0.5` second animations in parallel.

```
menuIntroDuration = duration menuIntro == 0.5
```

140 Of course, it is not possible to inspect every animation. In the following
 141 situation we have a custom operation `get`, the dual of `set` in the previous section,
 142 returning a `Float`. If the result of this value is used as the duration parameter,
 143 then we cannot know upfront how long this animation will last. Requesting to
 144 calculate the duration then results in a type error.

```

complicatedAnim = do v <- get; linearTo lens (For v) (To 10)
complicatedAnimDuration = duration complicatedAnim -- type error

```

145 Calculating a duration is a stepping stone towards other interesting features.
 146 One such example is sequentially composing animations with a relative offset.
 147 For example, to compose a first animation `anim1` with a second animation `anim2`
 148 which starts 0.5 seconds *before* the end of `anim1`.

```

relSeqAnim = relSequential anim1 anim2 (-0.5)

```

149 3.3 Expressiveness

150 In monadic DSLs the `>>=` and `return` combinators provide the needed expressiv-
 151 ity. When creating inspectable animations, `>>=` is a liability since it has limited
 152 inspectability. PaSe supports extension with custom control flow combinators.

153 The `onlyDone` animation shows all *done* items while hiding all to-do items.
 154 This could be implemented by first showing all items with the `showAll` anima-
 155 tion, since an item might have been hidden by a previous action, and then hiding
 156 all to-do items with the `hideToDo` animation. The definition is given below, while
 157 the definitions of `showAll` and `hideToDo` are omitted for brevity.

```

onlyDoneNaive = do showAll; hideToDo

```

158 However, we only intend to show completed items if needed. So instead we
 159 first check how many done items there are, if there are more than zero we play
 160 the previous version of `onlyDone`, otherwise we only hide the unfinished items.

```

onlyDone = do
  cond <- doneItemsGt0 -- check if more than 0 'done' items
  if cond then onlyDoneNaive else hideToDo

```

161 However, this formulation uses monadic features and is thus not inspectable.
 162 To make it inspectable, we utilize a custom combinator `ifThenElse`. We revisit
 163 this example in more detail in Section 5.

```

onlyDone = ifThenElse doneItemsGt0 onlyDoneNaive hideToDo

```

164 For this new combinator, we can define custom ways to inspect it. Since each
 165 branch might have a different duration, we do not choose to extract the duration
 166 but rather the *maximum* duration of the animation.

```

onlyDoneMaxDuration = maxDuration onlyDone -- = 1

```

167 Sections 2 and 3 gave a look and feel of the features of PaSe. In the following
 168 sections, we delve deeper into the internals of the implementation.

169 4 Implementation of PaSe

170 This section implements the previously introduced operations and redefines the
 171 animations to show the resulting type signature. We develop PaSe in the style of
 172 the `mtl` library¹⁰ which implements monadic effects using type classes [10]. This
 173 style is also called the finally tagless approach [3]. However, because the PaSe
 174 classes are not subclasses of `Monad`, they leave room for inspectability.

175 4.1 Specifying Basic Animations

176 The `mtl` library uses type classes to declare the basic operations of an effect.
 177 Similarly, we specify the `linearTo` operation using the `LinearTo` type class.

```
class LinearTo obj f where
  linearTo :: Traversal' s Float -> Duration -> Target -> f ()
```

178 The traditional `mtl` style would add a `Monad f` superclass constraint. As
 179 it hinders inspectability, we defer the addition of this constraint to the user.
 180 This allows the definition of animations which are, for example `Applicative`, if
 181 inspectability is needed or `Monad` if it is not.

182 The `linearTo` function is used to specify basic animations like `line1Out`,
 183 `line2In`, `menuSlideIn`, and `appFadeOut` from Section 2. As an example, we
 184 redefine `line1Out` with its type signature; the others are similar.

```
line1Out :: (LinearTo Application f) => f ()
line1Out = linearTo (navbar . underline1 . width) (For 0.25) (To 0)
```

185 4.2 Specifying Composed Animations

186 Section 2 used the combinators `sequential` and `parallel` for composing ani-
 187 mations. In this section, we describe these combinators in more detail.

188 **Sequential Composition** We reuse the `Functor-Applicative-Monad` hierar-
 189 chy for sequencing animations.

190 ehe `liftA2` function from the `Applicative` class, which has type
 191 `Applicative f => (a -> b -> c) -> f a -> f b -> f c`, takes two animations `f a`
 192 and `f b` and returns a new animation which plays them in order. The final result
 193 of the animation is of type `c`, which is obtained by using the function `a -> b -> c`
 194 and applying the results of the two played animations to it.

195 The `>>=` function from the `Monad` class, which has type `Monad f => f a ->`
 196 `(a -> f b)`, takes an animation `f a` and then feeds the result of this animation
 197 into the function `a -> f b` to play the animation `f b`.

198 The `sequential` function is a specialization of the `liftA2` function. It only
 199 applies to animations with a `()` return value, and trivially combines the results.

¹⁰ <http://hackage.haskell.org/package/mtl>


```

sequential :: (Applicative f) => f () -> f () -> f ()
sequential f1 f2 = liftA2 (\_ _ -> ()) f1 f2

```

200 Hence, the type signature for `selectBtn2Anim` contains an `(Applicative f)`
 201 constraint in addition to the `(LinearTo Application f)` constraint.

```

selectBtn2Anim :: (LinearTo Application f, Applicative f) => f ()
selectBtn2Anim = line1Out 'sequential' line2In

```

202 **Parallel Composition** We create our own `Parallel` type class for the `parallel`
 203 function¹¹. Its `liftP2` function has the same signature as `liftA2`, but the
 204 intended semantics of the `liftA2` implementation is parallel rather than se-
 205 quential composition. Technically they are interchangeable, but the relation of
 206 `Applicative` to `Monad` makes it more sensible for sequential composition seman-
 207 tics. The `parallel` function is a specialization of `liftP2`.

```

class Parallel f where
  liftP2 :: (a -> b -> c) -> f a -> f b -> f c

parallel :: (Parallel f) => f () -> f () -> f ()
parallel f1 f2 = liftP2 (\_ _ -> ()) f1 f2

```

208 With that in place we can give a type signature for `menuIntro`.

```

menuIntro :: (LinearTo Application f, Parallel f) => f ()
menuIntro = menuSlideIn 'parallel' appFadeOut

```

209 4.3 Running Animations

210 Now we create a new `Animation` data type that instantiates the above type
 211 classes to interpret PaSe programs as actual animations. We briefly summarize
 212 this implementation here and refer for more details to our codebase.¹²

213 The `Animation` data type, defined below, models an animation. It takes the
 214 current state `s` and the time elapsed since the previous frame. It produces a new
 215 state for the next frame, the remaining unused time and either the remainder of
 216 the animation or, if there is no remainder, the result of the animation. Note that
 217 the output is wrapped in a type constructor `m` to embed custom effects. We need
 218 the unused time when there is more time between frames than the animation
 219 uses. Then, the remaining time can be used to run the rest of the animation.

```

newtype Animation s m a = Animation { runAnimation ::
  s ->                                -- previous state
  Float ->                            -- time delta
  m ( s                                -- next state
    , Either (Animation s m a) a    -- remainder / result
    , Maybe Float )}                -- remaining delta time

```

¹¹ The `Alternative` class (https://en.wikibooks.org/wiki/Haskell/Alternative_and_MonadPlus) is not suitable as the laws are not the same.

¹² https://github.com/rubenpieters/anim_eff_dsl/tree/master/code

220 **LinearTo Instance** The `linearTo` implementation of `Animation` constructs
 221 the new state, calculates the remainder of the animation and the remaining
 222 delta time. The difference between the `linearTo` duration and the frame time
 223 determines whether there is a remaining `linearTo` animation or remaining time.

224 *Examples* We illustrate the behaviour on a tuple state `(Float, Float)`, of an
 225 `x` and `y` value. The `right` animation transforms the `x` value to 50 over 1 second.

```
right :: (LinearTo (Float, Float) f) => f ()
right = linearTo x (For 1) (To 50)
```

226 We run it for 0.5 seconds by applying it to the `runAnimation` function,
 227 together with the initial state (`s0 = (0,0)`) and the duration `0.5`. We instantiate
 228 the `m` type constructor inside `Animation` with `Identity` as no additional effects
 229 are needed; this means that the result can be unwrapped with `runIdentity`.

```
(s1, remAn1, remDel1) = runIdentity (runAnimation right s0 0.5)
-- s1 = (25.0, 0.0) | remAn1 = Left anim2 | remDel1 = Nothing
```

230 Running `right` for 0.5 seconds uses all available time and yields the new
 231 state `(25, 0)`. The remainder of the animation is the `right` animation with its
 232 duration reduced by `0.5`, or essentially `linearTo x (For 0.5) (To 50)`. Let
 233 us run this remainder for 1 second.

```
(s2, remAn2, remDel2) = runIdentity (runAnimation anim2 s1 1)
-- s2 = (50.0, 0.0) | remAn2 = Right () | remDel2 = Just 0.5
```

234 Now the final state is `(50, 0)` with result `()` and remaining time `0.5`.

235 **Monad Instance** For sequential animations we provide a `Monad` instance. Its
 236 `return` embeds the result `a` inside the `Animation` data type. The essence of
 237 the `f >=> k` case is straightforward: first, run the animation `f`, then pass its
 238 result to the continuation `k` and run that animation. We return the result of the
 239 animation, or, if there is an animation remainder, because the remaining time
 240 was used up, we return that remainder.

241 *Examples* Let us define an additional animation `up` which transforms the `y`
 242 value to 50 over a duration of 1 second. Additionally, we define an animation
 243 `rightThenUp` which composes the `right` and `up` animations in sequence.

```
up :: (LinearTo (Float, Float) f) => f ()
up = linearTo y (For 1) (To 50)

rightThenUp :: (LinearTo (Float, Float) f, Applicative f) => f ()
rightThenUp = right 'sequential' up
```

Running the `rightThenUp` animation for 0.5 seconds gives a similar result to running `right` for 0.5 seconds. We obtain the new state `(25, 0)`, an animation remainder `anim2` and there is no remaining time. Now the animation remainder is the rest of `rightThenUp`, which is half of `right` and `up`. So, when we run this animation remainder for 1 second, it will run the second half of `right` and the first half of `up`. This results in the state `(50, 25)`, the animation remainder `anim3` and no remaining delta time. This animation remainder is of course the second half of the `up` animation. If we continue to run that remainder, for example for 1 second, then we get the final state `(50, 50)` and the animation result `()`.

Parallel Instance The `liftP2` implementation runs the animations `f1` and `f2` on the starting state. We match on the cases where `f1` and `f2` finish with a result or an animation remainder and remaining time. We check which of the animations have finished and repackage them either into a result or a new remainder, using the result combination function where appropriate. When the longest of the two parallel animations is finished while not fully using the remaining delta time, we continue running the remainder of the animation.

Examples Let us run the animations `right` and `up` in parallel, which means that both the `x` and `y` value will increase simultaneously.

```
rightAndUp :: (LinearTo (Float, Float) f, Parallel f) => f ()
rightAndUp = right 'parallel' up
```

The result of running this animation for 0.5 seconds gives the state `(25, 25)` and no remaining time. If we continue the animation remainder we get the state `(50, 50)` and 0.5 seconds of remaining time.

4.4 Inspecting Animations

To inspect animations we instantiate them with `Const`. It wraps an `a` value and has a `b` phantom type parameter to trivially make it a functor.

```
newtype Const a b = Const { getConst :: a }
```

We might wonder why this extra work is necessary. After all, it is possible to obtain the duration of an animation by running the animation and keeping track of how long it takes. First, this is not an ideal approach for obtaining the duration. We might obtain erroneous results when doing this on conditional animations. Since only one branch of the conditional will be taken, while the other branch with a different duration might be taken in reality. Also, this approach is infeasible when there are effects embedded within the animation. Second, duration is one possible inspection target. Another example is tracking the used textures within an animation so they can be loaded automatically. For this to be possible we must run the inspection *before* the animation runs for the first time, since the textures must be loaded first.

279 **Inspecting LinearTo** To obtain the duration of a `linearTo` animation, we
 280 embed the duration in the `Const` wrapper.

```
instance LinearTo obj (Const Duration) where
  linearTo _ duration _ = Const duration
```

281 **Inspecting Applicative** It is not possible to inspect animations with a `Monad`
 282 constraint, but it is possible for animations with an `Applicative` constraint.
 283 The `Const` data type is not the culprit here, but rather the `>>=` method of the
 284 `Monad` class, which contains the limiting factor: a continuation function `a -> m b`.

285 **Inspecting Parallel** The duration of two parallel animations is the maximum
 286 of their durations. The `Par (Const Duration)` instance implements this.

```
instance Par (Const Duration) where
  liftP2 _ (Const x1) (Const x2) = Const (max x1 x2)
```

287 *Examples* The duration function is a specialization of the unwrapper function
 288 of the `Const` data type, namely `getConst`. We can feed our previously defined
 289 animations `selectBtn2Anim` and `menuIntro` from Section 2 to this function and
 290 obtain their durations as a result.

```
duration :: Const Duration a -> Duration
duration = getConst

selectBtn2AnimDuration :: Duration
selectBtn2AnimDuration = duration selectBtn2Anim -- = For 1.0

menuIntroDuration :: Duration
menuIntroDuration = duration menuIntro -- = For 0.5
```

291 When we try to retrieve the duration of a monadic animation, there is an
 292 error from the compiler: there is no `Monad` instance for `Const Duration`.

```
complicatedAnimDuration :: Duration
complicatedAnimDuration = duration complicatedAnim
-- No instance for (Monad (Const Duration))
```

293 4.5 Adding a Custom Operation

294 Custom operation are added by defining a corresponding class. For example, if
 295 we want to add a `set` operation, then we create the corresponding `Set` class.

```
class Set obj f where set :: Lens' obj a -> a -> f ()
```

296 Now, an animation using the `set` operation will incur a `Set` constraint.

```

checkIcon :: (Set CompleteIcon f, ...) => f ()
checkIcon = do ...; set (checkmark . color) green; ...

```

297 To inspect or run such an animation, we also need to provide instances for
 298 the `Animation` and `Const` data types. In the `Animation` instance, we alter the
 299 previous state by setting the value targeted by the `lens` to `a`. The duration of a
 300 `set` animation is 0, which is what is returned in the `Duration` instance.

```

instance (Applicative m) => Set obj (Animation obj m) where
  set lens a = Animation $ \obj t -> let
    newObj = Lens.set lens a obj
    in pure (newObj, Right (), Just t)

instance Set obj (Const Duration) where
  set _ _ = Const (For 0)

```

301 5 Interaction Between Inspectability and Expressivity

302 Haskell DSLs are typically monadic because the `>>=` combinator provides great
 303 expressive power. Yet, this power also hinders inspectability. This section shows
 304 how to balance expressiveness and inspectability with a custom combinator.
 305 This feature is *opt-in* in the sense that it is only required when inspectability
 306 is required. If that is no concern, then it is no problem to work with the `Monad`
 307 constraint.

308 Let us revisit the `onlyDone` animation from Section 3.3. The following defini-
 309 tion imposes a `Monad` constraint on `f`, making the animation non-inspectable.

```

onlyDone :: (LinearTo Application f, Get Application f,
  Set Application f, Monad f, Parallel f) => f ()
onlyDone = do
  cond <- doneItemsGt0
  if cond then onlyDoneNaive else hideNotDone

```

310 However, there is duration-related information we can extract. For example,
 311 the *maximum duration* is the largest duration of the two branches.

312 To express this idea in PaSe we introduce an explicit combinator to replace
 313 this particular use of `>>=`, namely an `if-then-else` construction.

```

class IfThenElse f where
  ifThenElse :: f Bool -> f a -> f a -> f a

```

314 This is similar to the `handle` combinator from the `DynamicIdiom` class [21]
 315 and the `ifS` combinator from the `Selective` class [18].

316 Now we can reformulate `onlyDone` in terms of this `ifThenElse` combinator¹³.

¹³ Using GHC's `RebindableSyntax` extension, it is possible to use the builtin
`if ... then ... else ...` syntax.

```

onlyDone :: (LinearTo Application f, Get Application f,
  Set Application f, Applicative f, Parallel f, IfThenElse f)
=> f ()
onlyDone = ifThenElse doneItemsGt0 onlyDoneNaive hideNotDone

```

317 We implement an appropriate `Animation` instance for `IfThenElse`.

```

instance (Monad f) => IfThenElse (Animation obj f) where
  ifThenElse fBool thenBranch elseBranch = do
    bool <- fBool
    if bool then thenBranch else elseBranch

```

318 Now, we can retrieve the maximum duration, using the `newtype` `MaxDuration`
 319 to signify this. The instance for `IfThenElse` retrieves the durations of the `then`
 320 and `else` branches and adds the greater value to the duration of the preceding
 321 animation inside the condition.

```

instance IfThenElse (Const MaxDuration) where
  ifThenElse (Const (MaxDur durCond)) (Const (MaxDur durThen))
    (Const (MaxDur durElse)) =
    Const (MaxDur (durCond + max durThen durElse))

```

322 This allows us to retrieve the maximum duration of the `onlyDone` animation.

```

onlyDoneMaxDuration :: MaxDuration
onlyDoneMaxDuration = maxDuration onlyDone -- = MaxDur 1.0

```

323 6 Interaction Between Callbacks and Inspectability

324 Many JavaScript animation libraries¹⁴ exist, most of which allow the user to add
 325 custom behavior (which the library has not foreseen) through callbacks. A good
 326 example is the GreenSock Animation Platform (GSAP), a widely recommended
 327 and mature JavaScript animation library with a variety of features.

328 6.1 Working with GSAP

329 `TweenMax` objects are the GSAP counterpart of the `linearTo` operation. Their
 330 arguments are similar: the object to change, the duration, and the target value
 331 for the property. For example, animation `right` moves `box1` to the right:

```

const right = new TweenMax("#box1", 1, { x: 50 });

```

332 We can add animations to a `TimeLineMax` to create a sequential animation.
 333 Below, we create `rightThenDown` which moves `box1` to the right and then down.

¹⁴ Examples: <https://greensock.com>, <https://animejs.com>, and <https://popmotion.io>.

```

const rightThenDown = new TimelineMax({ paused: true })
  .add(new TweenMax("#box1", 1, { x: 50 })))
  .add(new TweenMax("#box1", 1, { y: 50 }));

```

334 The `add` method takes the position on the timeline as an optional paramter.
 335 If we position both animations at point 0 on the timeline, they run in parallel.
 336 For example, the `both` animation below moves both `box1` and `box2` in parallel.

```

const both = new TimelineMax({ paused: true })
  .add(new TweenMax("#box1", 1, { x: 50 })), 0)
  .add(new TweenMax("#box2", 1, { x: 50 })), 0);

```

337 Timelines can also be embedded within other timelines.

```

const embedded = new TimelineMax({ paused: true })
  .add(both.play())
  .add(new TweenMax("#box1", 1, { y: 50 })), 0);

```

338 6.2 Callbacks and Inspectability

339 GSAP provides features related to inspectability. For example, we can use the
 340 `totalDuration` method to return the total duration of an animation. Ordinary
 341 animations correctly give their total duration when queried. For example, query-
 342 ing the duration of `embedded` correctly returns 2.

```

const embeddedDuration = embedded.totalDuration(); // = 2

```

343 However, if we want to provide animations similar to `onlyDone`, which con-
 344 tains an `if-then-else`, then the duration returned is not what we expect. The
 345 `add` method is overloaded and can also take a callback as parameter. Using the
 346 callback parameter we can embed arbitrary effects and control flow. For exam-
 347 ple, we can create a conditional animation `condAnim`, for which a duration of
 348 1 is returned. This is because any callbacks that are added to the timeline are
 349 considered to have duration 0, even if an animation is played in that callback.

350 The resulting duration of 1 is different from the expected total duration of
 351 the animation, which is 2. Of course, in general the duration of the animations in
 352 both branches could differ, which is what makes it difficult to provide a procedure
 353 for calculating the duration of an animation in this form.

```

const condAnim = new TimelineMax({ paused: true })
  .add(both.play())
  .add(() => { if (cond) { new TweenMax("#box1", 1, { x: 50 }) }
              else { new TweenMax("#box2", 1, { x: 50 }) } });
const condAnimDuration = condAnim.totalDuration() // = 1

```

6.3 Relevance of Duration in Other Features

A wrongly calculated duration becomes more problematic when another feature relies on this calculation. The *relative sequencing* feature needs the duration of the first animation, so the second animation can be added with the correct offset. For example, we can specify the position parameter `--0.5` to specify that it should start 0.5 seconds before the end of the previous animation.

```
const bothDelayed = new TimelineMax({ paused: true })
  .add(new TweenMax("#box1", 1, { x: 50 }), 0)
  .add(new TweenMax("#box2", 1, { x: 50 }), "--0.5");
```

This feature differs from ordinary sequencing such as with `sequential`. When we state that animation B must play 0.5 seconds before the end of animation A, then it is not possible to wait until animation A has finished to start running animation B. This is because animation B *should have started playing for 0.5 seconds already*. When we have the duration of animation A available, animation B can be appropriately scheduled.

This feature behaves somewhat unexpectedly when combined with a conditional animation. In the `relativeCond` animation below we add a basic animation followed by a conditional animation. Then we add an animation with a relative position. The result is that the relative position is calculated with respect to the duration of the animations before it, which was a duration of 1.

```
const relativeCond = new TimelineMax({ paused: true })
  .add(new TweenMax("#box1", 1, { x: 50 }), 0)
  .add(() => { if (cond) { new TweenMax("#box1", 1, { x: 100 });
    } else { new TweenMax("#box1", 1, { x: 0 }); } })
  .add(new TweenMax("#box2", 1, { x: 50 }), "--0.5");
```

Predicting the resulting behavior becomes much more complicated when conditional animations are embedded deep inside complex timelines and cause erroneous duration calculations. Clearly, being more explicit about control flow structures and their impact on inspectability like in PaSe helps providing a more predictable interaction between these features.

6.4 Relative Sequencing in PaSe

While not yet ideal from a usability perspective,¹⁵ PaSe does enable correctly specifying relative sequential compositions by means of `relSequential`.

```
relSequential :: forall c g.
  (c (Const Duration), c g, Applicative g, Delay g) =>
  (forall f. c f => f ()) -> g () -> Float -> g ()
relSequential anim1 anim2 offset = let
  dur = getDuration (duration anim1)
  in anim1 'sequential' (delay (dur + offset) *> anim2)
```

¹⁵ It requires `AllowAmbiguousTypes` (among other extensions) and explicitly instantiating the constraint `c` at the call-site.

379 Because this definition requires instances instantiated with `Const Duration`,
 380 it only works for animations whose duration can be analyzed. Now, we can
 381 correctly compose conditional animations sequentially using relative positioning.
 382 We use the `relMaxSequential` function to sequence animations with a maximum
 383 duration.

```
-- create synonym for multiple constraints
class (LinearTo Float f, IfThenElse f) => Combined f where
instance (LinearTo Float f, IfThenElse f) => Combined f where

relCond :: (LinearTo Float f, IfThenElse f, Applicative f) => f ()
relCond = relMaxSequential @Combined anim1 anim2 (-0.5)
```

384 7 Use Case

385 This section compares an implementation of a simplified Pac-Man game (Fig-
 386 ure 5) in Haskell with PaSe and TypeScript with GSAP both quantitatively and
 387 qualitatively. The quantitative evaluation compares development time and lines
 388 of code. The qualitative one compares different aspects of the libraries.




Fig. 5: Screenshot of the Pac-Man application.

389 7.1 Quantitative Evaluation

390 This section compares the PaSe and GSAP implementations on quantitative
 391 criteria. We consider the development time and lines of code for each module.

- 392 – **Development Time** The Haskell application was developed in ~ 1.5 work-
 393 ing days, while the TypeScript application took ~ 1 working day. We consider
 394 this approximately the same development time as the Haskell application was
 395 developed first, and thus contains design work shared by both applications.
 396 The developer is proficient in both languages.
- 397 – **Lines of Code (LOC)** Table 1 contains the LOC data (including whites-
 398 pace) for both applications. Their total LOCs are roughly the same. However,
 399 the Haskell code implements its own functionality for sprites and textures
 400 while we used the existing `Sprite` class of the `PixiJS` library in TypeScript.

Table 1: Lines of code comparison (including whitespace)

| Module | Haskell/PaSe (LOC) | % | TypeScript/GSAP (LOC) | % |
|--------------|--------------------|-------------|-----------------------|-------------|
| AnimDefs | 127 | 21% | 197 | 32% |
| Anims | 43 | 7% | 39 | 6% |
| Field | 48 | 8% | 77 | 12% |
| Game | 130 | 21% | 113 | 18% |
| Main | 36 | 6% | 23 | 4% |
| Sprite | 45 | 7% | / | / |
| Textures | 34 | 6% | 13 | 2% |
| Types | 10 | 2% | 3 | 0% |
| View | 139 | 23% | 158 | 25% |
| <i>Total</i> | <i>612</i> | <i>100%</i> | <i>623</i> | <i>100%</i> |

401 – **Relative LOC** Table 1 also contains the relative LOCs. The GSAP ani-
 402 mation definitions (AnimDefs) are slightly bigger because we had to embed
 403 effects in the animations due to differences in the used graphics library,
 404 and because of TypeScript’s relative verbosity. Using the timeline feature of
 405 GSAP, the code for simple animations is comparable to PaSe. However for
 406 more complex animations and those requiring embedded effects, there are
 407 more differences that we discuss in more detail in the qualitative evaluation.

408 7.2 Qualitative Evaluation

409 This section compares PaSe and GSAP on five qualitative criteria.

410 – **Eco-system** Animations are not created in isolation; they need to be cou-
 411 pled to a graphical backend to display them on the screen. GSAP’s maturity
 412 makes it a clear winner here. It is well integrated with the browser and sup-
 413 ports a rich set of features such as a variety of plugins, compatibility across
 414 browsers and support for animating a large range of DOM elements. Yet, for
 415 Pac-Man we only needed lenses for our own user-defined state.

416 – **Workflow** It is important that animations can be specified easily and con-
 417 cisely. Creating pure animations, without any embedded effects, are equally
 418 convenient in GSAP and in PaSe. However, more complex interactions with
 419 effects and control flow are simpler in PaSe. We saw this in the Pac-Man
 420 use case when implementing particle animations. A particle animation is an
 421 animation that creates an object, animates it and then destroys it again. We
 422 implemented a general wrapper for such animations which takes as input a
 423 function `Int -> Animation`, where the `Int` is the unique particle identifier,
 424 and a creation and deletion function for the particle. In the GSAP library
 425 we have to add the function to the timeline as a callback, which means its
 426 duration is considered to be `0`. This is problematic because the deletion of
 427 the particle should occur after its animation. This means that we are forced
 428 to manually calculate and provide the duration for the particle animation.

- 429 – **Performance** Both libraries perform equally acceptable on Pac-Man: no visible
430 glitches or lag at 60 frames per second (FPS) on an Intel core i7-6600U at
431 2.60 GHz with 8GB memory. We have also implemented a benchmark similar
432 to GSAP’s speed test¹⁶, which tests a large parallel animations. GSAP is
433 slightly more optimized currently as it handles 500 parallel animations at 60
434 FPS instead of PaSe’s 400. This could be remedied by further performance
435 improvements of PaSe, like fusing multiple parallel animations or improving
436 the **Animation** data structure, which are future work.
- 437 – **Extensibility & Inspectability** Extensibility and inspectability are key
438 features of PaSe. Both were useful for Pac-Man. Inspectability allowed extracting
439 all used textures in the animations to automate their loading. Extensibility
440 enabled the definition of the particle effect mentioned earlier. We
441 created a new **WithParticle** type class and implemented both an **Animation**
442 instance and a **Const** instance for the texture inspection. GSAP does not support
443 inspectability, and thus we did not implement the automatic loading of
444 textures. The particle animation function was implemented with callbacks
445 and implicit side-effects, which TypeScript allows anywhere.

446 8 Future Work

447 The main aspects to improve are currently related to the PaSe eco-system, such
448 as providing bindings to various graphics backends, extending the provided set
449 of features and improving the performance of the implementation.

450 Another avenue of future work is to explore trade-offs between the MTL style,
451 as used in this paper, or an initial encoding approach, as is typical in approaches
452 based on algebraic effects and handlers. The MTL style was chosen since it is
453 simpler presentation-wise, mainly on the extensibility aspect with regards to
454 different computation classes. However, we believe that implementation of some
455 features, such as the relative sequencing, is simpler in the initial approach.

456 An aspect not touched in this paper is *conflict management*. A conflict appears
457 when the same property is targeted by different animations in parallel. For
458 example, if we want to change a value both to 0 and 100 in parallel, what should
459 this animation look like? PaSe does no conflict management, and the animation
460 might look stuttery. GSAP, for example, resolves this by only enabling the
461 most recently added animation. However this strategy is not straightforwardly
462 mapped to the context of PaSe. Inspectability could provide a solution for this
463 problem by providing the possibility to detect conflicts.

464 9 Related Work

465 *Functional Reactive Programming* The origins of functional reactive programming
466 (FRP) lie in the creation of animations [4], and many later developments
467 use FRP as the basis for purely functional GUIs.

¹⁶ <https://greensock.com/js/speed.html>

PaSe focuses on easily describing *micro-animations*, which differ from general *animations* as considered by FRP. The latter can typically be described by a time-parameterized picture function `Time -> Picture`. While a subset of all possible animations, micro-animations are not easily described by such a function because many small micro-animations can be active at the same time and their timing depends on user interaction.

We have only supplied an implementation of PaSe on top of a traditional event-based framework, but it is interesting future work to investigate an implementation of the `linearTo`, `sequential` and `parallel` operations in terms of FRP behaviours and events.

Animation Frameworks Typical micro-animation libraries for web applications (with CSS or JavaScript) and animation constructions in game engines provide a variety of configurable pre-made operations while composing complex animations or integrating new types of operations is difficult. PaSe focuses on the creation of complex sequences of events while still providing the ability to embed new animation primitives. We have looked at GSAP as an example of such libraries and some of the limits in combining extensibility with callbacks and inspectability. PaSe is an exercise in improving this combination of features forward in a direction which is more predictable for the user.

Planning-Based Animations PaSe shares similarities with approaches which specify an animation as a plan which needs to be executed [12,17]. An animation is specified by a series of steps to be executed, the plan of the animation. The coordinator, which manages and advances the animations, is implemented as part of the hosting application. PaSe realizes these plan-based animations with only a few core principles and features the possibility of adding custom operations and inspection. A detailed comparison with these approaches is difficult, since their works are very light on details of the actual implementation aspect.

Inspectable DSLs Some DSLs for parsing [9,2,15], non-determinism [11], remote execution [5,6] and build systems [19] focus on inspectability aspects, yet none of them provide extensibility and expressiveness in addition to inspection.

10 Conclusion

We have presented PaSe, an extensible and inspectable DSL for micro-animations. PaSe focuses on compositional animations using sequential and parallel animations as basic building blocks. This is in contrast with other animation libraries typically focused on sequential composition and callbacks with implicit effects.

We utilized a to-do list application use case to explain the features of PaSe. In this use case we showed the additional features of PaSe: extensibility, inspectability and expressivity. We argue that the callback style of providing extensibility hurts the inspectability aspect of animations, which is found in for example the GreenSock Animation Platform. An implementation of the Pac-Man game confirms that this can be a problem even in simple applications.

References

1. Bederson, B.B., Boltman, A.: Does animation help users build mental maps of spatial information? In: INFOVIS 1999. pp. 28–35 (1999). <https://doi.org/10.1109/INFVIS.1999.801854>
2. Capriotti, P., Kaposi, A.: Free applicative functors. In: MSFP 2014. pp. 2–30 (2014). <https://doi.org/10.4204/EPTCS.153.2>
3. Carette, J., Kiselyov, O., Shan, C.: Finally tagless, partially evaluated: Tagless staged interpreters for simpler typed languages. *J. Funct. Program.* **19**(5), 509–543 (2009). <https://doi.org/10.1017/S0956796809007205>
4. Elliott, C., Hudak, P.: Functional reactive animation. In: ICFP 1997. pp. 263–273 (1997). <https://doi.org/10.1145/258948.258973>
5. Gibbons, J.: Free delivery (functional pearl). In: Haskell 2016. pp. 45–50 (2016). <https://doi.org/10.1145/2976002.2976005>
6. Gill, A., Sculthorpe, N., Dawson, J., Eskilson, A., Farmer, A., Grebe, M., Rosenbluth, J., Scott, R., Stanton, J.: The remote monad design pattern. In: Haskell 2015. pp. 59–70 (2015). <https://doi.org/10.1145/2804302.2804311>
7. Gonzalez, C.: Does animation in user interfaces improve decision making? In: CHI 1996. pp. 27–34 (1996). <https://doi.org/10.1145/238386.238396>
8. Heer, J., Robertson, G.G.: Animated transitions in statistical data graphics. *IEEE Trans. Vis. Comput. Graph.* **13**(6), 1240–1247 (2007). <https://doi.org/10.1109/TVCG.2007.70539>
9. Hughes, J.: Generalising monads to arrows. *Sci. Comput. Program.* **37**(1-3), 67–111 (2000). [https://doi.org/10.1016/S0167-6423\(99\)00023-4](https://doi.org/10.1016/S0167-6423(99)00023-4)
10. Jones, M.P.: Functional programming with overloading and higher-order polymorphism. In: Advanced Functional Programming, First International Spring School on Advanced Functional Programming Techniques, Båstad, Sweden, May 24-30, 1995, Tutorial Text. pp. 97–136 (1995). https://doi.org/10.1007/3-540-59451-5_4
11. Kiselyov, O.: Effects without monads: Non-determinism - back to the meta language. In: ML/OCaml 2017. pp. 15–40 (2017). <https://doi.org/10.4204/EPTCS.294.2>
12. Kurlander, D., Ling, D.T.: Planning-based control of interface animation. In: CHI 1995. pp. 472–479 (1995). <https://doi.org/10.1145/223904.223968>
13. van Laarhoven, T.: CPS-Based Functional References (2009), <https://www.twanvl.nl/blog/haskell/cps-functional-references>
14. Liang, S., Hudak, P., Jones, M.P.: Monad transformers and modular interpreters. In: POPL 1995. pp. 333–343 (1995). <https://doi.org/10.1145/199448.199528>
15. Lindley, S.: Algebraic effects and effect handlers for idioms and arrows. In: WGP 2014. pp. 47–58 (2014). <https://doi.org/10.1145/2633628.2633636>
16. McBride, C., Paterson, R.: Applicative programming with effects. *J. Funct. Program.* **18**(1), 1–13 (2008). <https://doi.org/10.1017/S0956796807006326>
17. Mirlacher, T., Palanque, P.A., Bernhaupt, R.: Engineering animations in user interfaces. In: EICS 2012. pp. 111–120 (2012). <https://doi.org/10.1145/2305484.2305504>
18. Mokhov, A., Lukyanov, G., Marlow, S., Dimino, J.: Selective applicative functors. *ICFP 2019* pp. 90:1–90:29 (2019). <https://doi.org/10.1145/3341694>
19. Mokhov, A., Mitchell, N., Peyton Jones, S.: Build systems à la carte. *PACMPL* **2**(ICFP 2018), 79:1–79:29 (2018). <https://doi.org/10.1145/3236774>
20. Wadler, P.: Comprehending monads. In: LFP 1990. pp. 61–78 (1990). <https://doi.org/10.1145/91556.91592>

- 558 21. Yallop, J.: Abstraction for web programming. Ph.D. thesis, University of Edin-
559 burgh, UK (2010)